

Generative AI

Exam Preparation Worksheet (2025/26)

Miguel Cabral Pinto & amigos

Made with the context of the entire GenAI course from MEI @ UC. Questions marked ★ are actual exam questions; support project-related claims with experimental observations. Questions shown in grey are optional — they cover niche detail unlikely to appear in a 2-hour exam; do the black ones first.

Contents

Model Family Quick Reference	3
1 Generative AI Foundations	4
2 Classical Generative Methods	5
3 Autoencoders	7
4 Variational Autoencoders	8
5 Generative Adversarial Networks	10
6 Diffusion Models	12
7 Evaluation	14
8 Latent Space Navigation	15
9 LLMs as Sequence Generators	17
10 Modern Image Generation	19
11 Project 1 — ArtBench Generative Modeling	21
12 Project 2 — Image-to-Prompt Inversion	22
Answer Key	23

Key Equations Reference

Variational Autoencoder

$$\mathcal{L}_{\text{ELBO}} = \underbrace{\mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)]}_{\text{reconstruction}} - \underbrace{\text{KL}(q_\phi(z|x) \parallel p(z))}_{\text{regularisation}} \quad (1)$$

$$z = \mu_\phi(x) + \sigma_\phi(x) \odot \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, I) \quad (\text{reparameterisation trick}) \quad (2)$$

KL decomposition for diagonal Gaussian posteriors:

$$-\frac{1}{2} \sum_j (\mu_j^2 + \sigma_j^2 - \log \sigma_j^2 - 1)$$

Generative Adversarial Networks

$$\mathcal{L}_{\text{GAN}} = \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] \quad (3)$$

$$\mathcal{L}_{\text{WGAN}} = \mathbb{E}[f(x)] - \mathbb{E}[f(G(z))], \quad f \text{ is 1-Lipschitz} \quad (4)$$

$$\mathcal{L}_{\text{GP}} = \lambda \mathbb{E}_{\hat{x}} [(\|\nabla_{\hat{x}} f(\hat{x})\|_2 - 1)^2], \quad \hat{x} = tx + (1-t)G(z) \quad (\text{WGAN-GP}) \quad (5)$$

DDPM

$$q(x_t | x_0) = \mathcal{N}(\sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t)I), \quad \bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s) \quad (\text{one-shot forward}) \quad (6)$$

$$\mathcal{L}_{\text{simple}} = \mathbb{E}_{t, x_0, \varepsilon} [\|\varepsilon - \varepsilon_\theta(x_t, t)\|^2] \quad (7)$$

$$\hat{\varepsilon}_{\text{cfg}} = \varepsilon_\theta(x_t, t, \emptyset) + w (\varepsilon_\theta(x_t, t, c) - \varepsilon_\theta(x_t, t, \emptyset)) \quad (\text{classifier-free guidance}) \quad (8)$$

Rectified Flow / Flow Matching

$$x_t = (1-t)x_0 + tx_1, \quad x_1 \sim \mathcal{N}(0, I) \quad (\text{straight-line interpolant}) \quad (9)$$

$$\mathcal{L}_{\text{RF}} = \mathbb{E}[\|v_\theta(x_t, t) - (x_1 - x_0)\|^2] \quad (10)$$

β -VAE

$$\mathcal{L}_{\beta\text{-VAE}} = \mathbb{E}[-\log p_\theta(x|z)] + \beta \cdot \text{KL}(q_\phi(z|x) \parallel p(z)), \quad \beta > 1$$

LoRA

$$W' = W + \Delta W = W + BA, \quad B \in \mathbb{R}^{d \times r}, \quad A \in \mathbb{R}^{r \times k}, \quad r \ll \min(d, k)$$

Temperature Sampling

$$P_\tau(x_t = w_j | x_{<t}) = \frac{\exp(z_j/\tau)}{\sum_k \exp(z_k/\tau)}$$

Model Family Quick Reference

	VAE	VQ-VAE	GAN	DDPM	Flow (RF)
Training obj.	ELBO (MSE + KL)	Recon + codebook + commit	Minimax / Wasserstein	$\ \varepsilon - \varepsilon_\theta\ ^2$	$\ v_\theta - (x_1 - x_0)\ ^2$
Latent space	Continuous, $\mathcal{N}(0, I)$	Discrete codebook	Implicit ($z \sim \mathcal{N}$)	Markov noisy chain	Straight ODE
Sample quality	Blurry	High (with prior)	Sharp	High	High
Mode coverage	Good	Good	Poor (collapse risk)	Excellent	Excellent
Speed	Fast	Fast + prior	Fast	Slow (1000 steps)	Few steps
Likelihood	Approx (ELBO)	None	None	Approx	None

1 Generative AI Foundations

Generative AI encompasses any system that constructs novel instances of a target domain. This section establishes the formal distinction between generative and discriminative models, surveys the four main paradigms (symbolic, statistical, evolutionary, connectionist), and examines the broader capabilities generative models provide beyond sample synthesis.

1.1 Which of the following best describes a generative model?

- A) A model that learns a decision boundary $P(Y | X)$ to classify inputs.
- B) A model that learns the data distribution $P(X)$ and can sample new instances from it.
- C) A model that retrieves the nearest stored example from a memory bank.
- D) A model that maps inputs to fixed-length feature vectors for downstream classifiers.

1.2 Evaluate the following claim and explain your reasoning.

“A generative model can always be used as a discriminative classifier via Bayes’ rule, but a discriminative model cannot be used to generate new samples.”

1.3 Explain why minimising reconstruction error on the training set alone is not sufficient as the sole objective for a generative model. What additional property must the learned distribution have, and how does this connect to the tension between memorisation and generalisation?

1.4 Describe the “generative paradox.” Give one concrete example of a model family that resolves each tension it describes.

1.5 Match each approach to its paradigm. Write the number in the blank.

Approach	Paradigm
a AARON (Harold Cohen)	_____
b Gaussian Mixture Models	_____
c L-systems	_____
d GANs	_____
e Genetic algorithms	_____
<i>Paradigms: (1) Statistical (2) Symbolic (3) Evolutionary / ALife (4) Connectionist / Deep</i>	

1.6 Explain the difference between “Generative AI” (broad) and “Generative ML” (narrow) as defined in the course. Give one example of a generative AI system that is not a generative ML system.

1.7 List three distinct capabilities that a generative model provides beyond sample synthesis (e.g., data augmentation). For each, give a concrete application domain.

2 Classical Generative Methods

Before deep learning, generative systems were built from explicit rules, evolutionary search, or statistical models over sequences. This section covers symbolic rule systems (AARON, shape grammars, L-systems), constraint-based methods (CSP/CLP), evolutionary approaches (genetic algorithms, biomorphs), and Markov-chain methods for music (EMI, Illiac Suite). Comparing paradigms across controllability, capacity for surprise, and scalability is a recurring exam theme.

2.1 Which statement best describes how AARON generates paintings?

- A) It trains on images and learns to reconstruct them via backpropagation.
- B) It applies stochastic rule selection from a hierarchical symbolic knowledge tree hand-crafted by Harold Cohen.
- C) It uses a genetic algorithm to evolve compositions toward a fitness function.
- D) It samples from a Markov chain trained on Cohen's past paintings.

2.2 Cohen's stated position throughout most of his career was: "AARON is creative; I am not the artist." Evaluate this claim.

2.3 L-systems apply all production rules simultaneously to every symbol in the current string, unlike context-free grammars which apply one rule at a time. Explain what class of natural phenomena this parallel rewriting naturally models, and give one example of a structure that a sequential grammar cannot faithfully represent but an L-system can.

2.4 Describe the generate-and-test approach used in the Illiac Suite. What are the two components and how do they interact?

2.5 Compare AARON (symbolic) and Evolved Virtual Creatures (evolutionary) on four dimensions: source of knowledge, degree of controllability, capacity for surprise, and role of the human designer.

2.6 Shannon's text experiments (1948) showed that increasing the order of a Markov model makes text more locally coherent. State one limitation of high-order Markov models that motivated the move to neural approaches.

2.7 Fill in the missing cells.

Dimension	Symbolic	Evolutionary	Statistical
Knowledge source	Human engineering	_____	Data (learning)
Controllability	High	_____	Medium
Novelty / Surprise	_____	High	Medium
Interpretability	_____	Low	Medium
Scalability	_____	Medium	Medium→High

2.8 (opt.) Explain the role of turtle graphics interpretation in L-systems. Give the axiom and two derivation steps of the algae example ($A \rightarrow AB$, $B \rightarrow A$, axiom A).

2.9 (*opt.*) In a Constraint Satisfaction Problem (CSP) used for generative purposes, what does the “generative space” equal?

- A) The set of all variable assignments tried before a solution is found.
- B) The set of all valid variable assignments satisfying all constraints.
- C) The search heuristic used by the solver.
- D) The number of constraints minus the number of variables.

2.10 ★ Harold Cohen’s AARON and evolutionary art systems (e.g., Karl Sims’ Evolved Virtual Creatures) represent two different philosophies for building generative systems. Compare these classical approaches in terms of the source of “creativity,” their capacity for surprise, and the role of the human designer.

3 Autoencoders

An autoencoder maps inputs through a bottleneck and reconstructs them, forcing the network to learn compact representations. This section covers the classical undercomplete autoencoder and its regularised variants: denoising (input corruption), sparse (activation constraints), contractive (Jacobian penalty), and the convolutional U-Net with skip connections for dense prediction.

3.1 What is the primary reason classical (undercomplete) autoencoders are not suitable as standalone generative models?

- A) They use MSE loss, which is not differentiable.
- B) Their latent space is discontinuous and unstructured, so sampling from it produces garbage.
- C) They cannot be trained with gradient descent.
- D) They always overfit to noise.

3.2 Explain why corrupting the input with noise and training to reconstruct the clean version teaches the network more than simply training to reconstruct clean input from itself. What does this training criterion encourage the model to learn about the structure of the data distribution?

3.3 Evaluate the following claim.

“A sparse autoencoder must have fewer hidden units than the input dimension.”

3.4 Explain the purpose of skip connections in a U-Net. In what sense does pooling “lose information,” and how do skip connections recover it?

3.5 Compare denoising, sparse, and contractive autoencoders on: (a) how regularisation is applied, (b) what the latent space looks like, and (c) one application each is well suited for.

3.6 Hinton & Salakhutdinov (2006) showed that deep autoencoders outperform PCA on 20Newsgroups for dimensionality reduction. What is the key structural advantage of a deep autoencoder over PCA?

- A) PCA cannot handle more than 100 dimensions.
- B) Autoencoders can learn nonlinear manifolds; PCA is restricted to linear projections.
- C) PCA requires labelled data; autoencoders do not.
- D) Autoencoders use cross-entropy loss; PCA uses Euclidean distance.

3.7 (opt.) Why are transposed convolutions sometimes replaced by nearest-neighbour upsampling followed by a regular convolution in decoder architectures? What artifact does transposed convolution introduce?

3.8 (opt.) Explain the Jacobian penalty in contractive autoencoders. What geometric property does it enforce on the mapping $f : x \mapsto z$?

4 Variational Autoencoders

Variational Autoencoders add a probabilistic prior to the latent space, converting the autoencoder into a proper generative model. This section derives the ELBO objective, explains why the marginal $p_\theta(x)$ is intractable, covers the reparameterisation trick, and examines key failure modes (posterior collapse) and extensions (β -VAE for disentanglement, VQ-VAE for discrete codes).

4.1 In a VAE, the encoder outputs:

- A) A single deterministic latent vector z .
- B) The parameters μ and σ of a Gaussian distribution over z .
- C) A discrete codebook index.
- D) A probability over all training images.

4.2 The ELBO has two terms. For each: (a) write what it computes, (b) describe what the model is encouraged to do if that term is maximised in isolation, and (c) describe the failure mode that results when only that term dominates training.

4.3 The VAE cannot maximise $\log p_\theta(x)$ directly because

$$p_\theta(x) = \int p_\theta(x|z) p(z) \, dz.$$

Explain why this integral is intractable.

4.4 What problem does the reparameterisation trick solve?

4.5 Describe “posterior collapse” in a VAE. Which term in the ELBO causes it, and what practical technique (with a one-sentence explanation) is used to avoid it?

4.6 Explain what β -VAE adds to the standard VAE objective and what disentanglement means geometrically. Why does increasing β improve disentanglement at the cost of reconstruction quality?

4.7 Evaluate the following claim.

“Sampling from $\mathcal{N}(0, I)$ and decoding always produces a realistic image in a well-trained classical autoencoder.”

4.8 VQ-VAE avoids posterior collapse because:

- A) It uses a higher learning rate than standard VAEs.
- B) The encoder must commit to a specific discrete codebook entry, so the decoder cannot ignore the latent code.
- C) It uses the Wasserstein distance instead of KL divergence.
- D) It trains the encoder and decoder separately.

4.9 Compare VAE and VQ-VAE on: (a) latent space type, (b) susceptibility to posterior collapse, (c) how new samples are generated, and (d) one downstream application enabled by each.

4.10 (*opt.*) Describe the three separate effects of the KL term on the encoder's posterior. What does each of the components $-\mu_j^2$, $-\sigma_j^2$, and $\log \sigma_j^2$ do individually?

4.11 (*opt.*) Explain the straight-through gradient estimator used in VQ-VAE. What problem does it solve and what approximation does it make?

4.12 ★ A Variational Autoencoder is trained by optimising a loss with two main terms. Explain the role of each term and what happens to the generated outputs when one dominates the other.

5 Generative Adversarial Networks

Generative Adversarial Networks frame generation as a two-player minimax game between a generator G and a discriminator D . This section covers the minimax objective and its connection to JS divergence, the main failure modes (mode collapse, vanishing gradients), and the major architectural developments: DCGAN (stable training rules), WGAN/WGAN-GP (Lipschitz constraints), CycleGAN (unpaired translation), and StyleGAN ($\mathcal{W}$ space and AdaIN).

5.1 Starting from the GAN minimax objective

$$\min_G \max_D \mathbb{E}[\log D(x)] + \mathbb{E}[\log(1 - D(G(z)))]$$

derive the optimal discriminator $D^*(x)$ for a fixed generator G . What value does $D^*(x)$ take when $p_g = p_{\text{data}}$? Explain why reaching this theoretical equilibrium is difficult in practice.

5.2 Explain why the original generator loss $\min_G \log(1 - D(G(z)))$ has vanishing gradients early in training. What is the non-saturating alternative and why does it help?

5.3 Minimising the GAN objective (with the optimal discriminator substituted in) is equivalent to minimising which divergence between p_{data} and p_g ?

- A) KL divergence
- B) Total variation distance
- C) Jensen–Shannon divergence
- D) Wasserstein-1 distance

5.4 Describe mode collapse precisely. Give one concrete example (e.g., MNIST). Then list three techniques that mitigate it.

5.5 DCGAN introduced specific architectural constraints that made GAN training stable for the first time on natural images. For each design choice below, explain the problem it solves and why the discarded alternative was problematic:

- Strided/fractional-strided convolutions instead of pooling
- Batch normalisation in both G and D
- Removal of fully connected layers
- LeakyReLU in D , ReLU in G

5.6 Evaluate the following claim.

“WGAN replaces the discriminator with a critic that outputs a probability between 0 and 1.”

5.7 Explain the Lipschitz constraint in WGAN and why it is necessary. Describe the difference between weight clipping (WGAN) and gradient penalty (WGAN-GP) as two ways of enforcing it.

5.8 Describe the CycleGAN setup. How many networks are involved? What is the cycle consistency loss and why does it prevent mode collapse in the unpaired translation setting?

5.9 Compare Pix2Pix and CycleGAN on: (a) data requirements (paired vs unpaired), (b) number of networks, (c) loss function components, and (d) a typical use case for each.

5.10 StyleGAN introduces a mapping network and an intermediate latent space $\mathcal{W}$. Explain what the mapping network does and why $\mathcal{W}$ is considered more disentangled than the input $\mathcal{Z}$ space. Why does this distinction matter when using StyleGAN for semantic image editing?

5.11 (*opt.*) In BigGAN, the “truncation trick” works by:

- A) Removing the top- k least likely samples at inference.
- B) Interpolating the latent input toward the mean, reducing variance and trading diversity for fidelity.
- C) Truncating the discriminator’s gradient updates to prevent overfitting.
- D) Quantising the latent space to a discrete codebook.

5.12 ★ Explain the adversarial training dynamic of GANs and describe mode collapse as a failure mode, including how you would detect it. You may reference the objective:

$$\mathcal{L}_{\text{GAN}} = \mathbb{E}[\log D(x)] + \mathbb{E}[\log(1 - D(G(z)))]$$

6 Diffusion Models

Diffusion models learn to reverse a gradual noising process, enabling high-quality generation through iterative denoising. This section covers DDPM’s forward and reverse processes, the noise-prediction training objective, DDIM as a deterministic fast-sampling alternative, the two main conditioning schemes (classifier guidance and classifier-free guidance), and latent diffusion for computational efficiency.

6.1 The DDPM forward process uses a noise schedule $\{\beta_t\}_{t=1}^T$ increasing from $\beta_1 \approx 0.0001$ to $\beta_T \approx 0.02$, which defines the one-shot noisy image:

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, I), \quad \bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

(a) Explain conceptually what the noise schedule controls and why β_t is monotonically increasing. (b) Why is the one-shot formulation computationally important even though the process is defined stepwise? (c) What constraint must $\bar{\alpha}_T$ satisfy for the reverse process to start from pure noise, and how does the schedule above achieve it?

6.2 Explain conceptually why predicting the noise ε added to x_0 (rather than predicting μ directly) is both mathematically equivalent and empirically more stable in DDPM.

6.3 Evaluate the following claim.

“DDPM training requires solving a minimax game between two networks, making it as unstable as GAN training.”

6.4 Describe the key difference between DDPM sampling and DDIM sampling. What property of DDIM enables (a) deterministic generation from a fixed noise map and (b) a 10–50× speedup at inference?

6.5 Explain classifier guidance (Dhariwal & Nichol, 2021). Write the guided noise prediction formula and explain what the guidance scale s controls.

6.6 Explain classifier-free guidance (CFG). What training trick enables the model to learn both conditional and unconditional denoising from a single network? Write the CFG formula and explain what happens when $w > 1$.

6.7 Compare classifier guidance and classifier-free guidance on: (a) extra components needed, (b) scalability to text conditioning, and (c) typical guidance scale values in practice.

6.8 Latent Diffusion Models (LDM) run the diffusion process in the latent space of a pretrained autoencoder. Explain the two-stage architecture, the approximate compression ratio for a 512×512 image, and why this reduces compute.

6.9 Imagen (Google Brain, 2022) found that the most important architectural choice for text–image alignment was:

- A) Using a U-Net with more parameters.
- B) Replacing the VAE with a VQ-VAE.
- C) Using a large language model (T5-XXL) as the text encoder.
- D) Training on a larger image dataset.

6.10 Compare VAE, GAN, and DDPM across: training stability, sample quality, mode coverage, generation speed, and whether an explicit likelihood is available.

6.11 (*opt.*) The U-Net used in DDPM is conditioned on the timestep t . How is t injected into the network?

- A) As an additional input channel concatenated to the noisy image.
- B) Via sinusoidal positional embeddings processed through FiLM conditioning layers.
- C) As a scalar multiplied by the final output layer's weights.
- D) Through cross-attention with a learned t -specific token.

6.12 ★ Describe the forward and reverse processes in DDPMs at a conceptual level. Then explain classifier-free guidance both in terms of mechanism and impact on image generation.

7 Evaluation

Evaluating generative models is non-trivial: quality is multi-dimensional and there is no single ground-truth metric. This section covers the main distributional metrics (FID, KID, IS), their underlying assumptions and failure modes, complementary metrics (Precision/Recall, LPIPS, CLIP-Score, CleanFID), and the experimental protocol choices that affect reproducibility and cross-paper comparability.

7.1 FID uses a pretrained Inception-v3 network to embed images into a feature space before measuring distributional distance. Explain why comparing distributions in feature space is necessary for meaningful evaluation of perceptual quality, rather than comparing in pixel space. What distributional assumption does FID then apply to those features, and what is one important scenario where this assumption breaks down?

7.2 Explain the Gaussian assumption embedded in FID. Name one scenario where this assumption is seriously violated and how it affects the validity of FID scores.

7.3 Evaluate the following claim.

“A model with a very low FID is necessarily diverse.”

7.4 Explain the conceptual difference between Precision and Recall for generative models (Kynkäänniemi et al.). Give an example of a failure mode that Precision catches but FID misses, and one that Recall catches but FID misses.

7.5 KID differs from FID in that:

- A) KID uses a different pretrained feature extractor.
- B) KID makes no Gaussian assumption and provides an unbiased estimator, making it more reliable at small sample counts.
- C) KID measures per-sample quality rather than distributional similarity.
- D) KID uses pixel space rather than feature space.

7.6 List three limitations of CLIPScore as a metric for text-to-image models. What metric would you pair it with for a more complete evaluation?

7.7 Explain the protocol design problem in evaluation. Name four experimental choices that, if reported differently, can make the same metric value scientifically meaningless.

7.8 Compare Inception Score (IS) and FID on: (a) whether they reference real data, (b) what each measures, and (c) one fundamental limitation of each.

7.9 (opt.) Explain the “CleanFID” problem. What is “implementation drift” and why can two papers report FID on the same model and reach contradictory conclusions?

8 Latent Space Navigation

The latent spaces of trained generative models contain structured semantic directions that can be found and traversed without retraining. This section covers interpolation methods (lerp, slerp), unsupervised direction discovery (GANSpace, SeFa), supervised direction discovery (InterFaceGAN), GAN inversion (Image2StyleGAN, pSp, e4e), and diffusion-era editing tools (Prompt-to-Prompt, Null-Text Inversion, ControlNet, DreamBooth).

8.1 When interpolating between two latent codes $z_0, z_1 \sim \mathcal{N}(0, I)$, linear interpolation $z(t) = (1 - t)z_0 + tz_1$ produces midpoints with lower expected norm than either endpoint. Explain why this is a problem for generation quality, and describe how spherical linear interpolation (slerp) avoids it geometrically.

8.2 Explain the fundamental reconstruction–editability trade-off in GAN inversion. Why does optimising aggressively for reconstruction push z^* into a region of low editability?

8.3 Explain how Prompt-to-Prompt edits images. What is “swapped” during the denoising process, and how does this preserve spatial layout while changing semantic content?

8.4 Evaluate the following claim.

“DDIM inversion works perfectly even with high classifier-free guidance weights.”

8.5 Explain the ControlNet “clone trick”: what is frozen, what is trainable, and why do zero convolutions guarantee that base model quality is never degraded during early training?

8.6 Compare Textual Inversion and DreamBooth on: (a) what parameters are updated, (b) training images required, (c) output file size, and (d) typical fidelity to the target concept.

8.7 Explain the two main families of approach for finding semantic editing directions in a GAN’s latent space: supervised (label-based) and unsupervised (data-driven). What information does each require, and what is the key trade-off between them? Give one example of each approach.

8.8 (opt.) GANSpace discovers editing directions by:

- A) Training a linear SVM on attribute labels in $\mathcal{W}$ space.
- B) Applying PCA to the intermediate activation space of the StyleGAN generator, finding principal axes of variation.
- C) Computing the eigenvectors of the first generator layer’s weight matrix $AA^\top$.
- D) Minimising a CLIP similarity loss over the latent space.

8.9 (opt.) Explain SeFa (Semantic Factorisation). What mathematical object is factorised, and why is this “architecture-intrinsic” rather than “empirically discovered”?

8.10 (*opt.*) Compare Image2StyleGAN, pSp, and e4e on: (a) inversion approach (optimisation vs encoder), (b) inference speed, (c) reconstruction quality, and (d) editability of the resulting code.

8.11 (*opt.*) Explain Null-Text Inversion. What two stages does it run, and what is the null embedding optimised to do?

9 LLMs as Sequence Generators

Large language models generate text autoregressively, one token at a time, by learning $P(x_t | x_{<t})$ over a large corpus. This section covers the autoregressive factorisation, subword tokenisation (BPE), decoding strategies and their quality–diversity trade-offs, prompting paradigms (zero-shot, few-shot, chain-of-thought), hallucinations, and the scaling laws that govern LLM performance.

9.1 The autoregressive factorisation $P(x_1, \dots, x_T) = \prod_{t=1}^T P(x_t | x_{<t})$ is an exact decomposition by the chain rule of probability. Explain why this makes generation sequential rather than parallel. Give one fundamental advantage and one fundamental disadvantage of sequential generation compared to a non-autoregressive model (e.g., a VAE or masked diffusion) that generates all tokens in one pass.

9.2 Top- p (nucleus) sampling selects the smallest set of tokens whose cumulative probability is $\geq p$. Its key advantage over top- k sampling is:

- A) It is faster to compute.
- B) The set size adapts dynamically: small when the model is confident, large when it is uncertain.
- C) It always selects exactly k tokens.
- D) It eliminates the need for temperature.

9.3 Describe the effect of temperature τ on the output distribution. What happens as $\tau \rightarrow 0$ and as $\tau \rightarrow \infty$? Give the formula for the temperature-scaled distribution.

9.4 Evaluate the following claim.

“Beam search always produces the globally highest-probability sequence.”

9.5 Explain the concept of “hallucination” in LLMs. Distinguish between factual hallucination and faithfulness hallucination. Why does the autoregressive training objective not explicitly discourage hallucination?

9.6 Explain why chain-of-thought (CoT) prompting improves performance on multi-step reasoning tasks. Describe the mechanism in terms of how intermediate tokens become part of the model’s context.

9.7 Compare zero-shot, few-shot, and in-context learning as control mechanisms for LLMs. For each, describe how the model’s behaviour is shaped and what assumption it relies on.

9.8 Explain how subword tokenisation (BPE) works at a high level and why it is preferred over word-level or character-level tokenisation for LLMs.

9.9 (opt.) The Chinchilla result (Hoffmann et al., 2022) showed that:

- A) Larger models always outperform smaller models regardless of data.
- B) For a fixed compute budget, many early LLMs were over-parameterised and under-trained on data, and performance is jointly determined by model size and dataset size.
- C) The optimal strategy is to train on as much data as possible at constant model size.
- D) Model size follows a power law with dataset size only.

9.10 ★ Autoregressive language models generate text one token at a time, and the sampling strategy strongly affects the output. Compare temperature scaling, top- k , and nucleus (top- p) sampling as three ways of controlling the balance between diversity and coherence.

10 Modern Image Generation

Recent advances have moved beyond the convolutional U-Net backbone and the DDPM training objective. This section covers transformer-based diffusion (DiT, MMDiT), flow matching and rectified flow as cleaner training objectives, single-step generation via distillation (consistency models, ADD), the FLUX/SD3 architecture family, fine-tuning methods (LoRA, DreamBooth), and the autoregressive revival enabled by better visual tokenisers.

10.1 DiT (Diffusion Transformer, Peebles & Xie 2023) differs from the U-Net backbone in that:

- A) It uses a completely different training loss.
- B) It replaces the convolutional encoder–decoder with a transformer operating on latent patches, enabling LLM-style scaling laws.
- C) It removes the VAE and works in pixel space.
- D) It uses adversarial training instead of denoising.

10.2 Explain rectified flow as a training objective. Write the interpolant, the target velocity, and the training loss. Why do straight-line trajectories allow fewer ODE integration steps at inference compared to DDPM?

10.3 Evaluate the following claim.

“DDPM is a special case of the flow matching framework.”

10.4 Explain consistency models. What is the “consistency property” $f_\theta(x_t, t) = x_0$ and how does it enable one-step generation? Distinguish between Consistency Training (CT) and Consistency Distillation (CD).

10.5 LoRA (Low-Rank Adaptation) works by:

- A) Replacing all weight matrices with their SVD decomposition.
- B) Freezing the base model and adding a trainable low-rank decomposition $\Delta W = BA$ alongside each weight matrix.
- C) Fine-tuning only the normalisation layers while freezing everything else.
- D) Distilling the base model into a smaller student network.

10.6 Explain the DreamBooth fine-tuning approach. What is the “class-preservation loss” and why is it needed?

10.7 (opt.) DiT uses adaLN-Zero conditioning to inject timestep and class information into each transformer block. Explain the mechanism: what is modulated, how is the conditioning signal used to produce the modulation, and what is the purpose of initialising the residual gate α to zero?

10.8 (*opt.*) Explain Adversarial Diffusion Distillation (ADD), used in SDXL-Turbo. How does combining a distillation loss with a GAN discriminator loss address the diversity gap left by standard distillation?

10.9 (*opt.*) Compare Stable Diffusion 3 / FLUX.1 and DALL-E 2 on: (a) architecture backbone, (b) training objective, (c) text encoder, and (d) conditioning mechanism.

10.10 (*opt.*) Explain the MMDiT (Multi-Modal Diffusion Transformer) block used in SD3 and FLUX. How does it differ from a standard DiT block in how text and image tokens interact?

10.11 (*opt.*) Compare the autoregressive revival (VAR, MAR, Emu3) against diffusion-based image generation on: (a) generation speed, (b) architectural unification with text models, (c) sample quality on benchmarks, and (d) the key enabler that made AR competitive again.

11 Project 1 — ArtBench Generative Modeling

Three model families compared on ArtBench-10 (32×32 paintings, 10 artistic styles): VAE, DCGAN (and a conditional variant), and DDPM (custom UNet and fine-tuned pretrained checkpoint). Evaluation: FID and KID computed over 5 000 generated vs. 5 000 real images, repeated across 10 seeds. Hyperparameter selection via Bayesian optimisation on a 20% training subset, then full retrain on 50 000 images.

11.1 ★ You trained VAE, DCGAN, and DDPM on ArtBench-10. Each model family has well-known theoretical properties that predict the kind of outputs it tends to produce. Discuss whether your experimental results confirmed or contradicted these theoretical expectations, and explain why. If one of your models performed surprisingly well or poorly on specific art styles, propose an explanation rooted in how that model family works.

11.2 ★ The project used FID and KID as the primary evaluation metrics. Critically evaluate this evaluation protocol: what aspects of generation quality do FID and KID capture well, what do they miss, and how could the protocol be misleading? If you could add one additional evaluation method to better understand the differences between your three models, what would it be and why?

11.3 ★ Suppose you are asked to repeat the project at 256×256 resolution instead of 32×32 . For each model family (VAE, DCGAN, DDPM), describe what would break or need to change architecturally. Which family benefits most from the increased resolution, and would you replace any of the three with a different approach at this scale? Justify your answer based on what you learned from the project.

11.4 On a 20% training subset, DCGAN achieves competitive FID scores; on the full 50 000-image dataset, a diffusion model achieves substantially lower FID. Explain this ranking reversal in terms of the data requirements of each training objective. What is each model's training loss actually optimising, and why does that determine how many examples are needed to generalise?

11.5 A conditional GAN trained with class labels consistently underperforms its unconditional DCGAN counterpart on FID, despite having access to additional information. Propose the most plausible failure mechanism, and describe what discriminator feature space visualisation (e.g. UMAP of penultimate-layer activations on real and generated samples) you would look for to confirm or reject your hypothesis.

11.6 The VAE uses a pixel-space MSE reconstruction loss. Explain why this specific objective predicts blurry outputs, and describe how switching to either a perceptual loss (VGG feature matching) or an adversarial loss would change the failure mode in a specific, predictable way.

11.7 (opt.) During Bayesian hyperparameter optimisation for a custom diffusion UNet at 32×32 resolution, fANOVA analysis reveals that the binary flag `use_attention` explains over 95% of FID variance, with all low-FID runs having attention *disabled*. (a) Propose an explanation for why bottleneck self-attention hurts at this resolution. (b) Would you expect the same finding at 256×256 ? Justify with reference to what self-attention adds at different spatial scales.

12 Project 2 — Image-to-Prompt Inversion

Task: given six images produced by a fixed Latent Consistency Model (LCM, 8 deterministic inference steps, seed encoded in filename), recover text prompts that re-render to visually similar images. Evaluation is image-side only: CLIP image-image similarity ($\uparrow$), LPIPS ($\downarrow$), pixel RMSE ($\downarrow$). Approaches explored: static captioners (CLIP-Interrogator, BLIP-2), gradient-based discrete optimisation (PEZ, LCM-Grad), and VLM-driven iterative refinement, with a population-based genetic algorithm as the submitted final method.

12.1 In image-to-prompt inversion, candidate prompts are judged by the images they produce when rendered by the generator, not by text similarity to a reference caption. Explain why image-side evaluation is essential for this task. Under what conditions does CLIP image-text cosine similarity (as a proxy) mislead you about actual render quality?

12.2 PEZ optimises soft token embeddings in CLIP space by gradient ascent on cosine similarity with a target image, then snaps to vocabulary tokens at evaluation time. Two projection schedules are possible: (a) project to discrete tokens only once at the end of optimisation, or (b) project at every gradient step with a straight-through estimator. Explain why end-only projection systematically fails and per-step projection recovers more useful prompts. What does this reveal about using a CLIP text-image proxy to drive generation quality?

12.3 ★ Suppose the image-to-prompt inversion task were repeated with a standard DDPM (1000 stochastic denoising steps) instead of an LCM (8 deterministic steps). For each of the following aspects of your approach, describe what would need to change and why: (a) the cost per prompt evaluation and its effect on the search budget; (b) the applicability of gradient-based methods that backpropagate through the denoiser; (c) the reproducibility of each evaluation under a fixed seed.

12.4 A VLM-driven iterative refinement loop conditions the language model on both the target image and the current LCM render at each step, then re-evaluates on the actual generator output. Explain how this closes the inversion loop compared to static captioning. What local-optimum failure mode can still occur in a single-chain search, and how does extending the approach to a population-based genetic algorithm mitigate it?

12.5 A continuous embedding oracle bypasses text entirely, optimising the full 77×768 cross-attention conditioning tensor directly against the generator’s denoising objective. It achieves better CLIP and LPIPS scores than any text-based method, yet its renders are visibly blurry. Explain both observations. What does the gap between the oracle’s CLIP score and the best text-projected version derived from that same oracle quantify about the fundamental limit of text-based prompt inversion?

12.6 (opt.) CLIP-Interrogator generates prompts in two stages: a BLIP captioner produces a short description of the image, then a curated phrase library is searched for entries whose CLIP embedding is most similar to the target image embedding. Explain why this achieves good CLIP image-image scores but limited LPIPS performance. What does this gap reveal about the relationship between semantic alignment in CLIP space and pixel-level perceptual reconstruction?

Answer Key

Section 1 — Generative AI Foundations

1.1 B.

1.2 True. A generative model $p(x, y)$ can recover $P(y|x) = p(x, y)/p(x)$ via Bayes' rule, so classification is always possible from a generative model. A discriminative model learns only $P(y|x)$ and has no representation of $p(x)$, so it cannot generate new x samples.

1.3 A model that memorises training examples exactly achieves zero reconstruction error but has learned the training set, not the underlying distribution — it cannot generate novel samples outside the data. A generative model must assign high probability to plausible inputs it has never seen. The tension is between fitting training data closely and generalising to the broader manifold. Each major family resolves this differently: VAEs use explicit KL regularisation to keep the latent space dense and samplable; GANs use an adversarial game that penalises any distribution mismatch; diffusion models use a noise-injection process that destroys and reconstructs structure, preventing trivial memorisation.

1.4 The generative paradox is the set of opposing tensions a generative model must navigate simultaneously: learn patterns without memorising data; be coherent without becoming repetitive; be diverse without losing structure; be realistic without overfitting; be novel without becoming random. VAEs resolve the memorisation/smoothness tension via KL regularisation; GANs address the realism/blur tension via adversarial loss; diffusion models handle the quality/diversity tension via iterative denoising.

1.5 a-2 (Symbolic), b-1 (Statistical), c-3 (Evolutionary/ALife), d-4 (Connectionist/Deep), e-3 (Evolutionary/ALife).

1.6 Generative AI (broad) includes any system that constructs novel artifacts from a target domain—rule-based systems, evolutionary algorithms, CSP solvers—none of which learn from data. Generative ML (narrow) specifically learns a distribution $p(x)$ from examples and samples from it. Example of non-ML generative AI: AARON, which generates paintings via hand-crafted symbolic rules without any data-driven learning.

1.7 (1) Anomaly detection via likelihood scoring—used in fraud detection. (2) Missing data imputation by sampling plausible completions—used in medical imaging. (3) Data augmentation to balance class distributions—used in rare-disease classification.

Section 2 — Classical Generative Methods

2.1 B.

2.2 False (contested). Cohen's original position was that he was the artist and AARON was a sophisticated tool. Over time he reconsidered and argued that AARON exhibited a form of creativity, since its outputs surprised even him. The question of where creativity "lives" remains unresolved; the answer depends on whether combinatorial novelty within a hand-coded rule space constitutes creativity.

2.3 Parallel rewriting models biological growth in which every cell divides simultaneously — no unit waits for others. This captures self-similar branching structures (plants, ferns, coral) where all branches elongate and bifurcate in lockstep at each growth step. A sequential grammar applying one rule at a time would produce asymmetric structures where some branches are always one derivation ahead of others, breaking the fractal self-similarity. For example, the algae model ($A \rightarrow AB$, $B \rightarrow A$) applied simultaneously gives Fibonacci-ratio string lengths; applying rules one at a time would break this ratio.

2.4 The Illiac Suite uses a generate-then-test loop: a stochastic process (random note generation) proposes candidates, and a rule-based screener applies counterpoint constraints to accept or reject each. Randomness provides exploration; rules enforce musical validity.

2.5

Dimension	AARON (Symbolic)	Evolved Virtual Creatures
Knowledge source	Hand-coded by Cohen over 40+ years	Emergent from fitness-driven search
Controllability	High (rules are explicit)	Low-medium (fitness shapes, not specifies)
Capacity for surprise	Low (combinations of known rules)	High (alien morphologies the designer never imagined)
Human designer role	Author of all aesthetic knowledge	Designer of fitness function and genotype representation

2.6 The state space grows exponentially with order n : a k -gram model over vocabulary V has V^{k-1} possible contexts, so most transitions are never observed in training data, causing data sparsity and unreliable estimates. RNNs learn a compressed hidden state rather than enumerating all contexts, removing this bottleneck.

2.7 Evolutionary column: Emergent (search); Low-Medium. Symbolic column: Low (novelty); High (interpretability); Low (scalability).

2.8 Turtle graphics assigns a physical interpretation to the L-system string: F moves forward drawing a line, + turns right, - turns left, [saves state,] restores state. Algae example: $n = 0$: A; $n = 1$: AB (apply $A \rightarrow AB$); $n = 2$: ABA (apply $A \rightarrow AB$ to A, $B \rightarrow A$ to B, simultaneously).

2.9 B.

2.10 AARON uses hierarchical production rules encoding all of Cohen’s aesthetic knowledge. Every output is a combination of pre-existing rules; novelty comes from the combinatorial space of rule interactions, not from learning or optimisation. The human designer is the sole source of all creativity, and no output can exceed the encoded knowledge. Evolutionary art systems (e.g., Sims) use a fitness function and genetic operators; no behavioural knowledge is hand-coded. Variation comes from mutation and crossover; fitness selection drives adaptation. The designer specifies the fitness landscape and representation, but the system produces morphologies and behaviours the designer never imagined. Key contrast: AARON’s creativity is bounded by Cohen’s knowledge; evolutionary systems generate novelty by search and can surprise their designers in ways AARON, by design, cannot.

Section 3 — Autoencoders

3.1 B.

3.2 Reconstructing clean input from itself has a trivial solution: the identity mapping. Corruption breaks this shortcut — the network must infer the clean signal from a degraded version, which requires learning what plausible clean data looks like rather than memorising a lookup. Specifically, the training criterion encourages the network to project noisy inputs back onto the manifold of valid clean inputs, learning local manifold geometry around each training point. This is also connected to score matching: the denoising criterion implicitly estimates the score $\nabla_x \log p(x)$, the gradient pointing from a noisy input toward higher-density regions of the data distribution. This is the same signal diffusion models exploit explicitly.

3.3 False. A sparse autoencoder is typically overcomplete ($\dim(z) \geq \dim(x)$). Sparsity is enforced not by reducing neurons but by constraining how many are active at once (via L1 penalty, KL divergence on activation rates, or top- k selection). This allows a richer vocabulary of features than a small bottleneck would.

3.4 Pooling progressively reduces spatial resolution, discarding fine-grained spatial detail in favour of higher-level features. Skip connections in U-Net concatenate the encoder’s feature maps (before pooling) directly to the corresponding decoder layers, giving the decoder access to high-resolution spatial information that pooling would otherwise destroy. This is essential for dense prediction tasks where output resolution must match input resolution.

3.5

	Denoising AE	Sparse AE	Contractive AE
Regularisation	Corrupts input; loss on clean reconstruction	Penalises active units (L1/KL/top- k)	Penalises Jacobian Frobenius norm
Latent space	Robust to noise perturbations	High-dimensional, only few dims active	Locally invariant; contracts input neighbourhoods
Application	Image/audio denoising; BERT pretraining	Interpretable feature learning	Robust representations for noisy sensors

3.6 B.

3.7 Transposed convolutions with kernel size not divisible by stride cause some output positions to receive contributions from more input values than others, producing a periodic “checkerboard” artifact. Replacing them with bilinear or nearest-neighbour upsampling followed by a regular convolution avoids this because the upsampling step has no learnable parameters and the subsequent convolution smooths the result uniformly.

3.8 The CAE adds a penalty $\lambda \|J_f(x)\|_F^2$ on the Frobenius norm of the encoder’s Jacobian $J_f = \partial f(x)/\partial x$. Small Frobenius norm means the encoder’s output changes very little when the input is perturbed, so the encoder contracts local volumes in input space to smaller volumes in latent space. This enforces local invariance: representations are stable to small nuisances while still discriminating between genuinely different inputs.

Section 4 — Variational Autoencoders

4.1 B.

4.2 Term 1 (reconstruction): $\mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)]$ pushes the decoder to faithfully reconstruct x from latent samples. Maximised alone, the encoder maps each input to a delta function (zero variance) to minimise sampling uncertainty, destroying the structured latent space and making generation by sampling from $\mathcal{N}(0, I)$ impossible — the decoder hits empty gaps. Term 2 (KL regularisation): $-\text{KL}(q_\phi(z|x)||p(z))$ forces the encoder posterior toward the prior, keeping the latent space dense and samplable. Maximised alone (or with $\beta \gg 1$), posterior collapse occurs: the encoder ignores input and maps everything to the prior, so the decoder outputs a weighted average regardless of z , producing blurry, identical samples.

4.3 The integral over $z \in \mathbb{R}^d$ is intractable because the decoder is a neural network with no closed-form marginal and d is typically hundreds of dimensions, so direct integration is computationally impossible. The reparameterisation trick solves the training problem: writing $z = \mu + \sigma \odot \varepsilon$ with $\varepsilon \sim \mathcal{N}(0, I)$ moves the stochastic node to ε , which has no learnable parameters, so gradients flow through μ and σ via backpropagation.

4.4 The sampling step $z \sim q_\phi(z|x)$ is stochastic and non-differentiable, so gradients cannot flow through it to update ϕ . The reparameterisation trick rewrites $z = \mu_\phi(x) + \sigma_\phi(x) \odot \varepsilon$ with $\varepsilon \sim \mathcal{N}(0, I)$; the randomness moves to ε (no parameters), and gradients now flow through μ_ϕ and σ_ϕ normally.

4.5 Posterior collapse occurs when the KL term dominates training, pushing the encoder’s posterior toward $\mathcal{N}(0, I)$ regardless of input so the decoder ignores z entirely and generates average outputs. KL annealing avoids this by starting training with the KL weight at zero and linearly increasing it over the first epochs, letting the encoder learn meaningful codes before regularisation pressure builds.

4.6 β -VAE multiplies the KL term by $\beta > 1$: $\mathcal{L} = \mathbb{E}[-\log p_\theta(x|z)] + \beta \cdot \text{KL}(q_\phi(z|x)||p(z))$. Disentanglement means each latent dimension captures a single independent generative factor. Increasing β pushes q_ϕ harder toward the isotropic Gaussian prior, encouraging independence between dimensions since the prior factorises. The cost is that the reconstruction term carries less weight, so the model cannot use as much latent capacity for reconstruction, producing blurrier outputs.

4.7 False. A classical autoencoder has no defined distribution over its latent space. Decoding a random $z \sim \mathcal{N}(0, I)$ typically produces garbage because the encoder may have mapped inputs to isolated points anywhere in $\mathbb{R}^d$, leaving large “holes” between them with no meaningful decoding.

4.8 B.

4.9

	VAE	VQ-VAE
Latent space	Continuous, $\mathcal{N}(0, I)$ prior	Discrete codebook, spatial grid of tokens
Posterior collapse	Susceptible with expressive decoder	Eliminated: encoder must commit to a discrete code
New samples	Sample $z \sim \mathcal{N}(0, I)$, decode	Need a learned prior (PixelCNN/Transformer), then decode
Downstream use	Smooth interpolation; drug molecule optimisation	DALL-E 1 image tokens; AudioLM speech tokens

4.10 $-\mu_j^2$ pulls latent means toward zero, preventing the encoder from mapping inputs to arbitrarily distant regions. $-\sigma_j^2$ prevents variance from exploding. $\log \sigma_j^2$ prevents variance from collapsing to zero, which would turn distributions into delta functions and eliminate the regularisation benefit.

4.11 The nearest-neighbour lookup $k^* = \arg \min_k \|z_e - e_k\|_2$ is non-differentiable (arg min has zero gradient everywhere). The straight-through estimator pretends in the backward pass that quantisation is an identity: gradients from the reconstruction loss pass directly from z_q back to z_e as if $z_q = z_e$. The approximation holds empirically after the codebook has stabilised and the quantisation error is small.

4.12 Term 1 (reconstruction): $\mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)]$ pushes the decoder to reconstruct the input accurately from samples of the encoder’s posterior. When this term dominates ($\beta \rightarrow 0$), the model reduces to a classical autoencoder with good reconstruction but an unstructured latent space, so sampling from $\mathcal{N}(0, I)$ hits gaps and produces noise. Term 2 (KL regularisation): $-\text{KL}(q_\phi(z|x)||p(z))$ pulls the encoder’s posterior toward the prior, making the latent space dense and samplable. When this term dominates ($\beta \gg 1$), posterior collapse occurs: the encoder ignores the input and maps everything to the prior, the decoder outputs the training-data mode, and all generated samples are identical.

Section 5 — Generative Adversarial Networks

5.1 For fixed G , maximising over D pointwise: each x contributes $p_{\text{data}}(x) \log D(x) + p_g(x) \log(1 - D(x))$ to the objective. Setting the derivative with respect to $D(x)$ to zero gives $D^*(x) = p_{\text{data}}(x)/(p_{\text{data}}(x) + p_g(x))$. When $p_g = p_{\text{data}}$, this equals 0.5 everywhere — the discriminator cannot distinguish real from fake. In practice, G and D are updated alternately with finite-capacity networks, so D is never truly optimal when G updates. If D becomes too strong before G improves, $D(G(z)) \approx 0$ and the generator loss $\log(1 - D(G(z)))$ saturates at zero gradient — one of the key instabilities the non-saturating loss and WGAN were designed to fix.

5.2 When D is confident $G(z)$ is fake, $D(G(z)) \approx 0$, so $\log(1 - D(G(z))) \approx 0$: the loss is flat with near-zero gradient. The non-saturating alternative has G maximise $\log D(G(z))$ instead: when $D(G(z)) \approx 0$ this equals $\log(0) \rightarrow -\infty$, giving a large gradient and a strong learning signal early in training. The fixed point of both objectives is the same; optimisation is substantially more stable with the non-saturating form.

5.3 C.

5.4 Mode collapse: the generator maps many different z values to the same output (or a small set of outputs) that reliably fools D , ignoring the training data’s diversity. On MNIST, a collapsed generator might produce only 3s and 7s. Three mitigation techniques: (1) Wasserstein loss with gradient penalty; (2) minibatch discrimination (discriminator sees groups of samples, penalising low-diversity batches); (3) feature matching (train G to match statistics of real data in D ’s feature space, not just fool D at the output level).

5.5 Strided convolutions: pooling discards spatial information with no learnable parameters; strided convolutions downsample while learning which spatial patterns to preserve, giving the discriminator better gradient signal about local structure. Batch normalisation: without normalisation, activations explode or collapse during the alternating adversarial updates, destabilising both networks; BN keeps each layer’s activations well-conditioned throughout training. No fully connected layers: FC layers do not respect spatial structure, wasting capacity on global correlations and making it harder to generate spatially coherent images; fully convolutional architectures exploit translation equivariance. LeakyReLU in D : standard ReLU sets negative activations to zero, killing gradients for those units and making the discriminator saturate; LeakyReLU ($\alpha \approx 0.2$) keeps a small gradient for negatives, preventing dead neurons in the discriminator which would stop the learning signal flowing to G .

5.6 False. WGAN’s critic outputs an unbounded real-valued score, not a probability between 0 and 1. The score difference between real and fake samples approximates the Wasserstein-1 distance via the Kantorovich–Rubinstein dual.

5.7 The Kantorovich–Rubinstein dual requires the critic f to be 1-Lipschitz: $|f(x) - f(y)| \leq \|x - y\|$ for all x, y , so that the dual formula holds exactly. Without this, the critic grows without bound. Weight clipping enforces Lipschitzness by clamping all critic weights to $[-c, c]$ after each update, but is crude: it can cause vanishing gradients if c is too small and distorts the critic’s capacity. WGAN-GP penalises the critic when its gradient norm deviates from 1, a smoother and more principled constraint that does not distort network capacity.

5.8 CycleGAN uses four networks: G_{AB} , G_{BA} , D_A , D_B . Cycle consistency requires $G_{BA}(G_{AB}(x)) \approx x$ and $G_{AB}(G_{BA}(y)) \approx y$. Without this constraint, each generator could map all inputs to a single output that fools its discriminator, since there are no paired examples to enforce correctness. The cycle constraint ties the two generators together, preventing either from ignoring input content.

5.9

	Pix2Pix	CycleGAN
Data	Paired (input–output pairs required)	Unpaired (two unrelated image sets)
Networks	2 (1 U-Net G , 1 PatchGAN D)	4 (G_{AB} , G_{BA} , D_A , D_B)
Loss components	Adversarial + L1 reconstruction	Adversarial ($\times 2$) + cycle consistency ($\times 2$)
Typical use case	Semantic map $\rightarrow$ photo; sketch $\rightarrow$ image	Horse $\leftrightarrow$ zebra; summer $\leftrightarrow$ winter

5.10 The mapping network is an 8-layer MLP that maps a sampled $z \in \mathcal{Z} \sim \mathcal{N}(0, I)$ to an intermediate code $w \in \mathcal{W}$. The $\mathcal{Z}$ space must match a fixed Gaussian prior, which forces entangled features together wherever the training distribution requires it. The mapping network can learn a non-linear transformation that unravels this entanglement: different dimensions of w more consistently control single semantic attributes (hair colour, pose, lighting) independently. For editing, this matters because directions in $\mathcal{W}$ more cleanly correspond to individual attributes — moving along a direction in $\mathcal{W}$ changes one feature while leaving others stable, whereas editing in $\mathcal{Z}$ tends to produce correlated, entangled changes.

5.11 B.

5.12 G and D play a minimax game: D maximises $\mathcal{L}_{\text{GAN}}$ by correctly labelling real inputs (output near 1) and fake inputs (output near 0); G minimises by producing samples $G(z)$ that D classifies as real. At Nash equilibrium, $D(x) = 0.5$ everywhere and $p_g = p_{\text{data}}$. Mode collapse: G learns to map many z values to the same output (or a small set) that reliably fools D , ignoring training data diversity. D then adapts to reject those outputs; G shifts to a new mode; the system cycles without converging. Detection: generate a large batch and observe lack of variety; apply a pretrained classifier to generated samples and check if only a few classes appear; monitor the variance of generated samples over training.

Section 6 — Diffusion Models

6.1 (a) The schedule controls how quickly signal is destroyed. Small early β_t values add tiny perturbations, allowing the U-Net to learn to remove fine details first; larger late values destroy coarse structure. Monotonically increasing β_t ensures a smooth signal-to-noise gradient from $t = 1$ to $t = T$ with no sudden jumps that would make the denoising task discontinuous. (b) The one-shot formula lets any training sample (x_t, t) be drawn directly from x_0 in a single operation, with no iteration. Training samples any timestep t uniformly and the target noise ε is sampled once — this is far more efficient than stepping through the chain one step at a time for each training example. (c) $\bar{\alpha}_T = \prod_{s=1}^T (1 - \beta_s) \approx 0$ must hold, so $x_T \approx \varepsilon \sim \mathcal{N}(0, I)$: virtually no signal from x_0 remains, making generation from pure Gaussian noise valid. With β_t increasing to 0.02 over $T = 1000$ steps, the cumulative product collapses to near zero.

6.2 Predicting ε and predicting μ_θ are mathematically equivalent since the reverse mean is a deterministic function of ε_θ and x_t . Predicting ε is empirically more stable because it is a zero-mean target regardless of timestep, giving the network a consistent regression objective. Predicting μ directly has a scale that changes with $\bar{\alpha}_t$, introducing instability across the 1000-step range.

6.3 False. DDPM training is a standard supervised regression problem: the loss is $\mathbb{E}\|\varepsilon - \varepsilon_\theta(x_t, t)\|^2$, a plain MSE with no adversarial game and no competing networks. This is one of the main reasons diffusion models displaced GANs for high-quality generation.

6.4 DDPM uses a Markovian stochastic process adding fresh noise at each step; all 1000 steps are required and outputs differ between runs. DDIM defines a deterministic ODE that shares the same training objective. (a) The same x_T always produces the same image (no stochastic z added), enabling inversion and interpolation in noise space. (b) Because the trajectory is deterministic and smooth, steps can be skipped—running only 50–100 of the 1000—without significant error, since the ODE solver can take larger steps on a smooth function.

6.5 Classifier guidance trains a separate classifier $p(y|x_t)$ on noisy images. At sampling:

$$\hat{\varepsilon}_{\text{guided}} = \varepsilon_\theta(x_t, t) - s\sqrt{1 - \bar{\alpha}_t} \nabla_{x_t} \log p(y | x_t)$$

The gradient pushes x_t toward images the classifier assigns high probability to class y . Scale $s = 0$ is unconditional; larger s increases class adherence at the cost of diversity.

6.6 During training, the condition c is randomly dropped (replaced with $\emptyset$, $\approx 10\%$ of batches), so a single network learns both $\varepsilon_\theta(x_t, t, c)$ and $\varepsilon_\theta(x_t, t, \emptyset)$. At sampling:

$$\hat{\varepsilon}_{\text{cfg}} = \varepsilon_\theta(x_t, t, \emptyset) + w(\varepsilon_\theta(x_t, t, c) - \varepsilon_\theta(x_t, t, \emptyset))$$

When $w = 1$ this is standard conditional generation. When $w > 1$, the guided prediction extrapolates away from the unconditional direction, amplifying the conditioning signal. Higher w improves prompt adherence and sharpness but reduces diversity and can produce oversaturated artifacts.

6.7

	Classifier guidance	Classifier-free guidance
Extra components	Separate noisy-image classifier	None (single network + condition dropout)
Text scalability	Poor (requires a text-image noisy classifier)	Excellent
Typical guidance scale	$s \approx 5\text{--}15$	$w \approx 3\text{--}8$

6.8 Stage 1: a pretrained VAE encodes a $512 \times 512 \times 3$ image to a $64 \times 64 \times 4$ latent ($\approx 48\times$ compression). Stage 2: the U-Net runs the DDPM denoising objective on this compact latent; at generation, the final denoised latent is decoded by the frozen VAE decoder. Compute scales roughly quadratically with spatial resolution, so operating at 64×64 instead of 512×512 gives $(512/64)^2 = 64\times$ fewer spatial tokens, dramatically reducing memory and compute while preserving perceptual quality through the VAE’s learned compression.

6.9 C.

6.10

	VAE	GAN	DDPM
Training stability	Good (ELBO, no game)	Poor (mode collapse, vanishing gradients)	Excellent (MSE regression)
Sample quality	Often blurry	Sharp, realistic	High quality + diverse
Mode coverage	Good	Mode collapse risk	Excellent
Speed	Fast (1 pass)	Fast (1 pass)	Slow (1000 steps; 50 with DDIM)
Explicit likelihood	Approx (ELBO lower bound)	None (implicit)	Approx (simplified ELBO)

6.11 B.

6.12 Forward process: a fixed Markov chain adds Gaussian noise over $T = 1000$ steps via schedule $\{\beta_t\}$; after T steps, the image is indistinguishable from $\mathcal{N}(0, I)$; no learning occurs here. Reverse process: a U-Net conditioned on timestep t learns to predict the noise ε added at each step; generation starts from

$x_T \sim \mathcal{N}(0, I)$ and iteratively denoises. CFG mechanism: during training, the condition c is randomly dropped ($\approx 10\%$ of batches), so the network learns both conditional and unconditional denoising from a single model. At sampling, $\hat{\varepsilon} = \varepsilon_\theta(\emptyset) + w(\varepsilon_\theta(c) - \varepsilon_\theta(\emptyset))$ extrapolates toward the conditional direction. Higher w improves prompt alignment and sharpness but reduces diversity and can oversaturate.

Section 7 — Evaluation

7.1 Pixel-space distances (e.g., MSE) do not track perceptual quality: a one-pixel translation yields high MSE yet looks identical, while a generated image with realistic textures has high pixel-MSE because it synthesises novel details rather than copying training pixels. Inception-v3 features are semantically meaningful — perceptually similar images map to nearby feature vectors regardless of exact pixel arrangement. FID then fits both real and generated feature sets as multivariate Gaussians and computes the Fréchet (Wasserstein-2) distance between them. The Gaussian assumption breaks down for multi-modal datasets (e.g., ArtBench-10 with 10 art styles): the true feature distribution has multiple separated clusters that a single Gaussian cannot represent, inflating the covariance and making FID values less interpretable.

7.2 FID models both real and generated feature distributions as multivariate Gaussians. Real image distributions are often multi-modal (images from different classes or art styles), which violates this assumption. With multiple modes, the fitted Gaussian has inflated covariance, making the Fréchet distance less interpretable and potentially rewarding a generated distribution that concentrates on one mode.

7.3 False. A model with severe mode collapse that produces one very high-quality mode can still achieve a low FID if that mode coincidentally matches the mean and covariance of the full real distribution. FID cannot separately measure diversity; Precision/Recall is needed to detect this failure.

7.4 Precision = fraction of generated samples that fall within the real data manifold (measures fidelity). Recall = fraction of the real manifold covered by generated samples (measures diversity). Precision catches mode collapse with sharp images: a model generates high-quality images from a small region, but recall is low—FID may miss this if the covered region matches the overall statistics. Recall catches a model that covers the full distribution with blurry outputs: precision is low, but the broad coverage may keep FID acceptable.

7.5 B.

7.6 (1) CLIPScore does not measure image quality or realism: a blurry or artifact-filled image can score highly if it matches the text description. (2) CLIP has biases from its training distribution and fails on fine-grained compositional reasoning (“a red ball to the left of a blue cube”). (3) It does not penalise spatial errors if the global embedding is similar. Pair with FID to capture distributional quality and diversity.

7.7 Four choices: (1) the image resizing algorithm applied before feature extraction; (2) the version and layer of the feature extractor (Inception-v3 vs CLIP ViT); (3) the number of samples used (FID variance is high below $\approx 10k$ images); (4) the random seed and sampling hyperparameters (temperature, guidance scale) used during generation.

7.8

	Inception Score (IS)	FID
References real data?	No (only classifies generated images)	Yes (compared to real distribution)
Measures	Quality (classifier confidence) + diversity (marginal entropy)	Distance between real and generated distributions in Inception feature space
Fundamental limitation	Can be gamed by training on ImageNet classes; no notion of distributional match	Assumes Gaussians; sensitive to sample count and preprocessing

7.9 FID depends on the image resize algorithm (bilinear, bicubic, Lanczos), JPEG compression decisions, and normalisation choices before feature extraction. Different codebases make different low-level choices,

shifting feature distributions and thus FID scores, even on identical images. A model could appear to “improve” FID simply by matching the preprocessing of the reference dataset rather than generating better images.

7.10 What FID and KID capture: both compare real and generated image distributions in Inception feature space, measuring whether the generated distribution matches the real distribution in mean, covariance (FID), and kernel-based statistics (KID); they are sensitive to both quality and diversity. What they miss: (1) neither measures prompt faithfulness; (2) neither measures per-image artifacts not captured by Inception features; (3) FID assumes Gaussian distributions; (4) neither captures fine-grained stylistic differences between art styles. How the protocol can be misleading: high CFG scales improve FID by sharpening samples at the cost of diversity; different preprocessing choices produce incomparable numbers; small sample sets inflate KID variance. Additional method: Precision and Recall (Kynkäänniemi et al.) separates fidelity (Precision) from diversity (Recall), making mode collapse visible independently of overall distributional distance.

Section 8 — Latent Space Navigation

8.1 Points sampled from $\mathcal{N}(0, I)$ in high dimensions concentrate near a hypersphere of radius $\sqrt{d}$ (by concentration of measure). The linear midpoint $z(0.5) = (z_0 + z_1)/2$ has expected norm $\approx \|z_0\|/\sqrt{2}$ for near-orthogonal high-dimensional vectors, placing it well inside the sphere in a low-density region the decoder was never trained on — this produces degraded or incoherent outputs. Slerp follows the great-circle arc on the hypersphere between z_0 and z_1 :

$$\text{slerp}(z_0, z_1; t) = \frac{\sin((1-t)\Omega)}{\sin \Omega} z_0 + \frac{\sin(t\Omega)}{\sin \Omega} z_1, \quad \Omega = \arccos\left(\frac{z_0 \cdot z_1}{\|z_0\| \|z_1\|}\right)$$

This keeps the norm constant throughout the path and stays in high-density prior regions, yielding semantically smooth intermediate samples.

8.2 The $\mathcal{W}$ space of StyleGAN corresponds to codes the training distribution actually used. Optimising the reconstruction loss aggressively pushes z^* into $\mathcal{W}^+$ (per-layer space) to fit fine-grained image details not expressible by the lower-capacity $\mathcal{W}$. Codes far from the training distribution’s mean were never encountered during training, so moving along a semantic direction from such a code produces unpredictable changes rather than the clean semantic edits the direction was calibrated for.

8.3 Prompt-to-Prompt injects the cross-attention maps from the source generation into the corresponding denoising steps of the target generation. The attention map encodes which spatial positions are controlled by which text tokens. Keeping the source maps fixed while changing the text tokens preserves structural layout (each position still “looks at” the same text token position) while changing the semantic content rendered at those positions.

8.4 False. CFG blends conditional and unconditional noise predictions with weight w . The DDIM inversion derivation assumes only the conditional score, so the inversion trajectory diverges from the forward trajectory when $w > 1$. At standard guidance weights ($w \approx 7.5$), accumulated error is large enough that decoding the inverted noise does not faithfully reconstruct the original image.

8.5 ControlNet clones the encoder half of the U-Net (trainable) while keeping the original U-Net frozen. Zero-initialised 1×1 convolutions connect the clone’s output to the frozen decoder. At training step 0, the zero convolutions multiply all clone outputs by zero, so the output is mathematically identical to the frozen base model regardless of what the clone receives. Base model quality can therefore never be degraded during early training, and the gradient landscape starts from the already-trained state.

8.6

	Textual Inversion	DreamBooth
Parameters updated	New token embedding S^* only (~kilobytes)	Full U-Net weights (or LoRA delta)
Training images	3–5	3–5
Output file size	~1 KB	Gigabytes (full) or 50–200 MB (LoRA)
Fidelity	Moderate; style and high-level features	High; identity preserved in fine detail

8.7 Supervised (label-based): train a linear classifier (e.g., SVM) on latent codes paired with binary attribute labels obtained from a pretrained attribute classifier. The decision boundary’s normal vector is the editing direction. Requires labelled attributes but produces semantically targeted, interpretable directions (e.g., “add glasses”, “change age”). Example: InterFaceGAN. Unsupervised (data-driven): discover principal axes of variation by applying PCA to latent codes or intermediate activations, requiring no labels. Directions correspond to the most dominant visual variations but may not map to human-interpretable attributes. Example: GANSpace. Trade-off: supervised gives control over specific named attributes at the cost of needing labelled data; unsupervised gives cheap, label-free exploration of dominant variation axes but with less interpretable semantics.

8.8 B.

8.9 SeFa computes the eigenvectors of $AA^\top$, where A is the weight matrix of the first generator layer. The principal eigenvectors are directions in $\mathcal{W}$ that produce the largest changes in early features, hence the largest semantic variation. This requires no samples and no labels: directions are derived directly from the weight matrix, making the method architecture-intrinsic.

8.10

	Image2StyleGAN	pSp	e4e
Approach	Per-image optimisation in $\mathcal{W}^+$	Encoder (feature pyramid)	Encoder + editability regularisation
Speed	Slow ($\sim 1\text{--}2$ min per image)	Fast (single forward pass)	Fast
Reconstruction quality	High	Good	Slightly lower
Editability	Low (deep $\mathcal{W}^+$ codes)	Moderate	High (regularised toward $\mathcal{W}$)

8.11 Stage 1 (pivot trajectory): standard DDIM inversion on the source image, producing reference noisy latents $\{x_T^*, \dots, x_1^*\}$. Stage 2 (null-text optimisation): at each timestep t , the unconditional embedding $\emptyset_t$ is optimised (model weights frozen) so that CFG-guided denoising from x_t^* exactly reproduces x_{t-1}^* from the pivot. This resolves the CFG-induced drift while preserving editing capability via the unchanged source prompt.

Section 9 — LLMs as Sequence Generators

9.1 The chain rule decomposition is exact for any distribution — no independence assumptions are made. However, each conditional $P(x_t | x_{<t})$ depends on all prior tokens, so x_t cannot be sampled until x_{t-1} is known, making parallelism during generation impossible; each step requires a full forward pass through the network. Advantage: the model can condition on arbitrarily long context to generate each token, capturing long-range dependencies exactly with no approximation in the factorisation — this is standard maximum-likelihood training. Disadvantage: generation latency scales linearly with sequence length ($O(T)$ forward passes vs. $O(1)$ for a VAE), which becomes a bottleneck for long sequences or real-time applications; non-autoregressive models can generate all tokens simultaneously but must make conditional independence approximations that reduce coherence.

9.2 B.

9.3 Temperature τ scales logits before softmax: $P_\tau(x_t = w_j | x_{<t}) = \exp(z_j/\tau) / \sum_k \exp(z_k/\tau)$. As $\tau \rightarrow 0$, the distribution peaks at the highest-logit token (greedy decoding, fully deterministic). As $\tau \rightarrow \infty$, the distribution becomes uniform. In practice, $\tau \approx 0.7\text{--}0.9$ for creative tasks and $\tau \approx 0.1\text{--}0.3$ for factual tasks.

9.4 False. Beam search maintains the top- k partial sequences and expands greedily—it is a heuristic best-first search, not an exhaustive search. It is not guaranteed to find the globally highest-probability sequence; in practice it tends to find overly conservative, repetitive outputs.

9.5 Hallucination is confident generation of content that is false or unsupported by context. Factual hallucination: invented facts (non-existent citations, wrong historical dates). Faithfulness hallucination: contradicting information explicitly present in the provided context. The autoregressive training objective rewards next-token prediction correctness, not factual accuracy: truth is not represented in the training signal, so the model has no direct incentive to be accurate.

9.6 With CoT, the model generates intermediate reasoning tokens that become part of the context $x_{<t}$ for subsequent tokens. Each reasoning step can use preceding steps as working memory, allowing multi-step inference that would exceed the model’s effective context compression if attempted in a single token. The mechanism relies on the model having been trained on text that includes step-by-step reasoning, so the CoT prompt activates the corresponding continuation distribution.

9.7 Zero-shot: task specified by instruction alone; relies on general pre-trained knowledge and instruction-following ability. Few-shot: 2–10 input–output examples prepended in the prompt; the model infers the task from the pattern; relies on in-context learning. In-context learning (general): any context—examples, demonstrations, instructions—adapts the model’s continuation without weight updates; relies on the model having seen similar patterns during pretraining.

9.8 BPE starts with individual characters and iteratively merges the most frequent adjacent pair in the corpus into a new compound symbol, until the vocabulary reaches a target size. The result is variable-length subwords: common words become single tokens; rare words decompose into shorter pieces. This avoids the out-of-vocabulary problem of word-level tokenisation while producing shorter sequences than character-level tokenisation.

9.9 B.

9.10 Temperature: divides logits by τ before softmax. $\tau < 1$ sharpens the distribution (more coherent, more deterministic); $\tau > 1$ flattens it (more diverse, less coherent). It does not constrain which tokens are eligible. Top- k : restricts sampling to the k most probable tokens, eliminating the long tail. The problem is that k is fixed regardless of distribution shape—sometimes 10 tokens are natural candidates, sometimes 1000 are. Nucleus (top- p): selects the smallest set of tokens whose cumulative probability reaches p (≈ 0.9 – 0.95). The set adapts: small when the model is confident (one token has near- p probability), large when uncertain. This makes the diversity–coherence balance input-dependent, which is why it is the current default. In practice, top- p and temperature are combined: temperature shapes the distribution first; top- p then cuts the remaining tail.

Section 10 — Modern Image Generation

10.1 B.

10.2 Rectified flow interpolates linearly: $x_t = (1 - t)x_0 + tx_1$, $x_1 \sim \mathcal{N}(0, I)$. The target velocity for this interpolant is $u_t = x_1 - x_0$, a constant vector per sample pair. Training loss: $\mathbb{E}\|v_\theta(x_t, t) - (x_1 - x_0)\|^2$. DDPM paths are curved because the stochastic forward process introduces variance at each step, requiring many small corrective integration steps. Rectified flow paths are straight lines, so the ODE can be integrated with larger steps (or fewer steps with an Euler solver) without significant approximation error.

10.3 True. DDPM corresponds to choosing a stochastic probability path with a specific variance schedule within the flow matching framework. Choosing the DDPM variance schedule as the probability path recovers the DDPM training objective; choosing a linear interpolant recovers rectified flow. Flow matching is strictly more general.

10.4 The consistency property requires $f_\theta(x_t, t) = x_0$ for every t along a given diffusion trajectory. One-step generation: start from $x_T \sim \mathcal{N}(0, I)$ and apply $f_\theta(x_T, T)$ directly to obtain x_0 . Consistency Training (CT) trains from scratch with a self-consistency loss between adjacent trajectory points: $f_\theta(x_t, t)$ must match $f_\theta(x_{t'}, t')$ for nearby t, t' , using a stop-gradient target network. Consistency Distillation (CD) uses a pre-trained multi-step diffusion model as teacher, training f_θ to match the teacher’s two-step output in one step; this produces higher quality than CT at the same step count.

10.5 B.

10.6 DreamBooth fine-tunes all U-Net weights on 3–5 reference images of a specific subject, binding the subject to a rare identifier token (e.g., “a [V] dog”). Without additional constraints, fine-tuning causes catastrophic forgetting: the model becomes highly specialised and loses the ability to generate other subjects or backgrounds. The class-preservation loss simultaneously fine-tunes on synthetic images from the original model for the same broad class (e.g., “a dog”), penalising deviation from the prior generation and preserving the model’s general capability.

10.7 adaLN-Zero uses the conditioning signal (timestep t and class label) as input to a small MLP that outputs scale γ and shift β parameters for each transformer block’s LayerNorm. The normalised activations are modulated as $\gamma \odot \text{LN}(x) + \beta$ before each sublayer (attention and MLP). An additional scalar

α , initialised to zero, multiplies the sublayer output before adding to the residual: $x \leftarrow x + \alpha \cdot \text{sublayer}(x)$. With $\alpha = 0$ at initialisation, every block is an identity function regardless of what the sublayer computes — the model starts as a trivially initialised network that produces outputs entirely determined by the embedding, and blocks are “switched on” incrementally as α grows during training. This guarantees stable gradient flow from the first update.

10.8 ADD combines two losses: (1) a distillation loss matching the student’s 1-step output to the multi-step teacher’s output distribution; (2) a GAN discriminator loss penalising the student when its outputs look perceptually unrealistic compared to real images. Standard distillation alone produces blurry averaged outputs because the student minimises a distribution-level loss that rewards the “average” of multiple plausible outputs. The discriminator provides a direct perceptual feedback signal that penalises this averaging and rewards individual-sample realism.

10.9

	SD3 / FLUX.1	DALL-E 2
Architecture	MMDiT (transformer on latent patches, two token streams)	CLIP prior + cascaded U-Net diffusion decoder
Training objective	Rectified flow	DDPM
Text encoder	CLIP-L + CLIP-G + T5-XXL (concatenated)	CLIP text encoder only
Conditioning	Joint attention in MMDiT (text and image tokens attend mutually)	Cross-attention between U-Net features and CLIP embeddings

10.10 A standard DiT block uses one-directional cross-attention: image patch tokens attend over text token embeddings (separate Q from image, K and V from text). MMDiT maintains two separate token streams—one for image patches, one for text—each with its own Q, K, V projections. Inside each block, all tokens form a single sequence for the attention operation, so every image token attends to every text token and vice versa. This bidirectional full joint attention in every layer gives the text conditioning deeper influence on image generation, which substantially improves text rendering and compositional accuracy.

10.11

	Autoregressive (VAR, MAR, Emu3)	Diffusion (DDPM, LDM, FLUX)
Speed	VAR $\sim 10\times$ faster than DiT at matched quality (parallel per scale level)	Slow (20–50 steps); 1-step via distillation recovers speed
Architectural unification	Native: same transformer, vocabulary, training loop as LLMs	Separate architecture; text via cross-attention adapters
Quality on benchmarks	VAR: 1.80 FID on ImageNet 256	DiT-XL/2: 2.27 FID; LDM-based leads on T2I benchmarks
Key enabler	Better visual tokenisers (MAGVIT-v2, FSQ) + next-scale prediction	Latent diffusion made high-resolution generation tractable

10.12 VAE at 256×256 : the encoder–decoder architecture scales (add layers), but the MSE reconstruction loss amplifies blurriness at high resolution since the model spreads probability mass over many equally plausible high-frequency pixel configurations. A perceptual loss and/or shifting to LDM (encode to compact latent first, diffuse there) would be necessary. DCGAN at 256×256 : training instability worsens severely because the discriminator can exploit high-frequency artifacts to identify fakes; generating sharp 256×256 images in a single pass is extremely hard. The standard fixes are progressive growing (ProGAN) or the StyleGAN architecture with its mapping network and style injection. DDPM at 256×256 : direct pixel-space diffusion is prohibitively expensive—U-Net self-attention scales quadratically with spatial dimension ($256^2 = 65,536$ positions). The necessary change is Latent Diffusion: encode $256 \times 256 \times 3$ to roughly $32 \times 32 \times 4$ with a frozen VAE ($\approx 64\times$ compression), then run DDPM in that compact latent. Most benefit: DDPM/LDM benefits most from higher resolution. At 32×32 the quality differences

between models are compressed; at 256×256 DDPM’s superior mode coverage and fine-detail generation pull decisively away from VAE and GAN. Would you replace architectures? VAE $\rightarrow$ VQ-VAE-2 for sharper textures; DCGAN $\rightarrow$ StyleGAN2 for stable training; plain DDPM $\rightarrow$ LDM is not a replacement but an essential wrapper. A defensible decision would be to drop GANs entirely at this scale, using LDM as the primary model and LoRA for efficient fine-tuning.

Section 11 — Project 1: ArtBench Generative Modeling

11.1 VAE: theory predicts blurry outputs due to the MSE reconstruction loss averaging over multiple plausible high-frequency configurations; this should be confirmed by low IS and high FID/KID. DCGAN: theory predicts mode collapse risk and training instability, but sharper outputs than VAE when stable; expect moderate FID and visible texture. DDPM: theory predicts the best sample diversity and quality given enough data, via a stable MSE noise-prediction loss over a well-defined forward process; expect dominant FID on the full dataset but weaker performance on small subsets. A common surprising result is DCGAN outperforming DDPM on a 20% subset: this contradicts the expectation that diffusion always wins and is explained by data efficiency — DDPM must estimate the score $\nabla_x \log q(x_t)$ across the full data manifold, which requires dense coverage that 10 000 images cannot provide. A GAN discriminator can overfit to a smaller real distribution faster and still provide useful gradients to the generator. If a particular art style (e.g. abstract expressionism) yields consistently worse FID than photorealistic styles, the explanation is higher intra-class variance: the model must cover a wider distribution per class, making distributional matching harder for all families.

11.2 FID and KID compare real and generated image distributions in Inception-v3 feature space (after the final pooling layer). They capture: overall distributional fidelity (does the generated set match the real set in feature space?), diversity (mode collapse shows as a narrow feature cluster far from the real cloud), and coarse perceptual quality. What they miss: (1) within-sample coherence — a model generating plausible textures at each location independently could score well without global structure; (2) style specificity — Inception features encode object identity for ImageNet, not artistic style; (3) human perceptual preference and composition; (4) multi-modal distributions — FID fits a single Gaussian to features, which breaks down for ArtBench’s 10-style mixture. The protocol can be misleading when a model is pretrained on a different distribution: fine-tuning from CIFAR-10 can produce visually coherent samples that are systematically shifted from ArtBench features, inflating FID beyond what visual inspection suggests (the features encode CIFAR-10 priors, not ArtBench statistics). A useful addition: Inception Score (IS), which measures both objectness ($p(y|x)$ low entropy) and diversity ($p(y)$ high entropy) on generated images alone, independent of the real data distribution. This directly separates fidelity failures from diversity failures, complementing FID/KID which conflate both.

11.3 DCGAN’s discriminator is trained with a binary real/fake signal on the available real images. With 10 000 images the discriminator has enough signal to identify plausible vs. implausible texture statistics, driving the generator toward locally convincing outputs. DDPM’s training objective — predicting added noise $\hat{\epsilon}_\theta(x_t, t)$ — requires the UNet to learn the full conditional score $\nabla_{x_t} \log q(x_t | x_0)$ across all noise levels and across the entire data manifold. Estimating this score accurately requires dense coverage of the manifold: with 10 000 samples the score is well-estimated near training points but unreliable elsewhere, so the reverse process can leave the manifold and produce incoherent images. With 50 000 samples the manifold is covered sufficiently for the score to guide the reverse process reliably from any starting noise. The GAN’s discriminator does not need to model the full density — it only needs to separate real from the current generator’s fake distribution, which is a comparatively easier target and is achievable with fewer real examples.

11.4 The most plausible mechanism is class-wise mode collapse: the cGAN generator converges to a per-class fixed point (a single representative image per class label) that reliably fools the discriminator without covering intra-class diversity. This occurs because the discriminator receives the class label as an additional input channel, giving it a shortcut — it learns class-specific cues to distinguish real from fake, which produces gradients pushing the generator toward the single most discriminator-fooling image per class rather than exploring the full intra-class distribution. Discriminator feature space evidence: a UMAP of penultimate-layer activations for real and generated images should show tight, degenerate fake clusters for each class that do not overlap with the spread of real clusters. By contrast, a well-generalising cGAN would show fake cluster KDE boundaries spreading across and mixing with their corresponding real clusters, indicating the generator covers the full intra-class diversity.

11.5 MSE minimises $\mathbb{E}[\|x - \hat{x}\|^2]$, which is minimised by the conditional mean of the output distribution

given the input. When multiple high-frequency textures (different brush strokes, fine details, edge orientations) are equally plausible reconstructions, the decoder predicts their average, washing out texture and contrast. Adding a perceptual loss (VGG feature matching) does not eliminate averaging but shifts the objective to higher-level feature channels — the model averages in perceptual feature space rather than pixel space, which tends to preserve edges and mid-frequency structure while still blurring fine-scale texture. Replacing with an adversarial loss changes the failure mode qualitatively: the decoder must now produce outputs that fool a discriminator, incentivising sharp, specific textures. However, this introduces mode collapse risk — the model may produce sharp but low-diversity outputs, trading blurriness for reduced coverage of the data distribution.

11.6 (a) At 32×32 , the UNet bottleneck operates on a $4 \times 4 = 16$ -position feature map. At this spatial scale, the receptive fields of residual convolution blocks already span the full image, so global context is implicitly available without attention. Inserting a self-attention block adds parameters and optimization complexity without providing any representational benefit the convolutions cannot already provide. The attention module may also interfere with the gradient flow from the timestep embedding, destabilising training. (b) At 256×256 , the equivalent bottleneck would operate on a 32×32 or larger feature map (1024+ positions). Residual convolutions have fixed local receptive fields and cannot efficiently aggregate information across spatially distant regions. Self-attention is necessary for global coherence — for example, ensuring that the color of a painted sky is consistent across the full width of a 256×256 image. The expected finding would reverse: low-FID runs would have attention *enabled*, which is exactly the design choice in standard DDPM and LDM UNets at high resolution.

Section 12 — Project 2: Image-to-Prompt Inversion

12.1 Text-to-image generators are not invertible: many distinct prompts can produce similar images, and conversely a small change in wording can cause a large change in the render. CLIP image-text cosine similarity measures alignment between a text embedding and an image embedding, but the generator uses only the text embedding as conditioning — two prompts with very similar CLIP text embeddings can produce very different renders if the generator responds differently to small embedding differences. CLIP image-text similarity misleads when: (1) a prompt accurately describes the target semantically but triggers a different composition or colour palette in the generator (high CLIP text-image score, high LPIPS); (2) stylistic modifiers that strongly affect the render (“Julie Dillon style”, “photorealistic”, “anime”) have low literal text-image similarity but dominate the generator output; (3) a vague but fluent prompt has moderate CLIP text-image similarity to many images but generates nothing close to the target.

12.2 End-only projection: gradients flow through continuous embeddings that are never constrained to the vocabulary manifold, so the optimiser moves them arbitrarily far from any real token. At the end, snapping to the nearest vocabulary token loses most of the optimised signal because real tokens are sparsely distributed in the embedding space — the optimal continuous point may map to a meaningless or unrelated token. The continuous objective can reach high cosine similarity while the corresponding realizable prompt scores near zero in render quality. Per-step projection: at each step the embeddings are snapped to vocabulary before the next gradient step (via a straight-through estimator that passes gradients as if the projection did not happen). This keeps the optimised prompt within the feasible set at all times, ensuring the CLIP score meaningfully reflects a realizable prompt. The fundamental tension exposed: even a per-step projected prompt with high image-text CLIP may produce a render with low image-image CLIP, because the generator’s decoder introduces structure that CLIP does not capture in the text-image alignment score.

12.3 (a) Standard DDPM with 1000 steps requires 1000 UNet forward passes per render vs. 8 for LCM — roughly 125× the compute per prompt evaluation. A search method that evaluates many candidates (e.g. a GA with population 8 over 8 generations runs ≈ 56 renders) becomes nearly infeasible at this cost. Either the search budget must shrink dramatically, or render feedback must be replaced by a cheaper proxy, which is less accurate. (b) Gradient-based methods (PEZ, LCM-Grad) backpropagate through the denoiser. In a standard DDPM reverse step, $x_{t-1} = \mu_\theta(x_t, t) + \sigma_t z$ where $z \sim \mathcal{N}(0, I)$ is freshly sampled. Gradients cannot flow through independent noise samples without explicit reparameterisation of the full 1000-step chain, making backpropagation through the reverse process intractable in practice. LCM’s consistency mapping is deterministic (no per-step noise injection), which is precisely why backpropagating through the denoising objective as a surrogate loss is tractable. (c) LCM renders are deterministic under a fixed seed and schedule: the same prompt always produces the same image, so repeated evaluations are identical. DDPM’s stochastic reverse process samples fresh noise at every step, meaning the same prompt

can produce slightly different images across calls unless the complete noise sequence is explicitly fixed and reproduced. Reproducibility under DDPM requires recording and replaying the full noise sequence for every evaluation.

12.4 Static captioners (CI, BLIP-2) produce a prompt from the target image once, with no knowledge of what the generator will actually render from that prompt. The loop is open: mismatches between the caption’s description and the generator’s interpretation are invisible and uncorrectable. The VLM refinement loop closes the loop by rendering the current prompt, observing both the target and the render, and asking the VLM to propose a correction conditioned on the visible discrepancy. The feedback signal is actual render quality, not an embedding proxy. Single-chain local optimum: if the initial prompt is already near-optimal (e.g. CLIP-Interrogator already produces the highest achievable CLIP score for a given target), the VLM may generate only small surface variations and never escape the local neighbourhood. All seeds then collapse to the same prompt from iteration zero, and the chain is stuck. A genetic algorithm maintains a population of 8 prompts, each potentially in a different region of the render-quality landscape. Crossover operators recombine prompts from two parents, allowing the search to jump between local optima. Elitism preserves the best individuals across generations, ensuring the population does not regress.

12.5 Blurriness: the oracle minimises the UNet’s expected denoising error $\mathbb{E}_t \|\hat{\epsilon}_\theta(x_t, \tau) - \epsilon\|^2$ over uniformly sampled timesteps and noise realisations. The minimiser of an expected squared error is the conditional mean, so the oracle converges to a conditioning that produces the average of all plausible renders rather than any single sharp realisation. High-noise timesteps (large t) dominate the uniform average and constrain only low-frequency structure, so the oracle never commits to fine-grained pixel detail. Better CLIP despite blurriness: CLIP embeds semantic content (subject identity, global palette, scene type) and is largely insensitive to high-frequency detail — a correctly identified subject in the right scene achieves high CLIP score whether the render is sharp or blurry. Better LPIPS despite blurriness: LPIPS with AlexNet features responds primarily to perceptual structure (edges, textures at mid frequency), which the oracle’s accurate layout still captures. Oracle-to-text gap: the oracle achieves mean CLIP 0.917; the best text-projected version derived from the same oracle achieves ≈ 0.853 . This gap quantifies the information lost when the optimal conditioning tensor is forced back into the discrete vocabulary — the best realizable text prompt is strictly worse than the optimal embedding. This is a lower bound on the cost of text discretisation and explains why text-based methods cannot fully close the inversion gap even when given the optimal continuous starting point.

12.6 CLIP-Interrogator proceeds in two stages. Stage 1: a BLIP captioner generates a short natural-language description of the target image (e.g. “a glass of orange juice with slices”). Stage 2: a curated phrase library covering styles, artists, modifiers, and media descriptors is searched for the entries whose CLIP ViT-L/14 image embedding is most similar to the target image embedding; the top-scoring phrases are appended to the caption to form a Stable Diffusion-style tag prompt. This achieves good CLIP image-image scores because both phrase selection and final evaluation use the same CLIP encoder: the phrases are chosen to maximise similarity in CLIP space, and the render inherits whatever stylistic content those phrases evoke in the generator. Limited LPIPS performance results from the mismatch between CLIP space and pixel space: CLIP is invariant to many structural and colour details that are perceptually important. A prompt that maximises CLIP semantic alignment may produce a render that matches the subject and global palette but differs substantially in composition, lighting angle, and spatial layout — all of which LPIPS penalises. The gap reveals that semantic similarity in CLIP space is necessary but not sufficient for pixel-level perceptual reconstruction.