

- ① ② ROC curves are related to the discriminative power of a feature by plotting sensitivity against specificity. A good feature separates classes well alone. We can calculate the area under the curve for each feature's ROC curve and select the best ones.
- ③ That feature is a perfect classifier and there is a threshold that perfectly separates the classes.
- ④ Features 2 and 3. Since the AUCs are similar, we should select the least redundant feature (f_2 and f_3 are the least correlated).

⑤

⑥ Some examples:

- Reduce dimensionality with LDA and then use MDC for classification.
- Train binary classifiers and decompose the multiclass problem as one of the following problems:
 - One-Against-All :
train one classifier per class (that class against every other class). Ideally, only 1 $d_i(x)$ should be positive while all other $d_j(x), j \neq i$ should be negative, then, choose class w_i .
 - One-Against-One :
train $k(k-1)/2$ binary classifiers, each $d_{ij}(x)$ is trained to discriminate between w_i and w_j ($d_{ij}(x) > 0$ if x belongs to w_i). The final result comes from majority voting.

→ Error-Correcting Output Codes (ECC) is the general framework (QAA and QAO are instances of it). Here we have a "coding matrix" that determines which classes each binary classifier discriminates. The final decision comes from comparing the vector of binary classifier outputs with the code words for each class (e.g., with Hamming distance metric) to find closest match.

⑥ $D=3$ (3 rows in each mean vector)

⑦ $K=3$ (3 mean vectors)

⑧ 1- Get projected mean for each class

$$\mu_K = W^T u_K$$

$$\mu_1 = \begin{pmatrix} 0.26 & -0.77 & -0.58 \end{pmatrix} \begin{pmatrix} +0.93 \\ -0.96 \\ -0.89 \end{pmatrix} \approx \begin{pmatrix} 1.01 \\ -0.03 \end{pmatrix}$$

$$\mu_2 = W^T \begin{pmatrix} -0.26 \\ 0.19 \\ -0.27 \end{pmatrix} \approx \begin{pmatrix} 0.24 \\ 0.04 \end{pmatrix}$$

$$\mu_3 \approx \begin{pmatrix} -1.24 \\ -0.01 \end{pmatrix}$$

2- Project sample

$$x' = W^T x \Rightarrow W \begin{pmatrix} 2 \\ 2 \\ 3 \end{pmatrix} = \begin{pmatrix} -2.76 \\ 0.16 \end{pmatrix}$$

3- Compute scores

$$d^2 = [(x_1 - \mu_K)^2 + (x_2 - \mu_{K_2})^2]^2$$

$$w_1 = [(-2.76 - 1.01)^2 + (0.16 + 0.03)^2]^2 \approx 14.2$$

w_2 : Do the same for the others and select
 w_3 : the best score

(iii) .

MDC decision boundary:

$$(m_1 - m_2)^T x - 0,5 (\|m_1\|^2 - \|m_2\|^2) = 0$$

in LDA space:
 $x \rightarrow W^T x$
 $m_k \rightarrow u'_k$

$$(u'_1 - u'_2)^T (W^T x) - 0,5 (\|u'_1\|^2 - \|u'_2\|^2) = 0$$

$$(AB)^T = B^T A^T$$

$$\underbrace{(W(u'_1 - u'_2))^T x}_{\text{weights vector}} - \underbrace{0,5 (\|u'_1\|^2 - \|u'_2\|^2)}_{\text{bias}} = 0$$

$$w = W(u'_1 - u'_2) = \begin{pmatrix} 0,26 & -0,75 \\ -0,77 & 0,65 \\ 0,58 & 0,12 \end{pmatrix} \begin{pmatrix} 0,77 \\ -0,07 \end{pmatrix} = \begin{pmatrix} -0,25 \\ -0,64 \\ -0,46 \end{pmatrix}$$

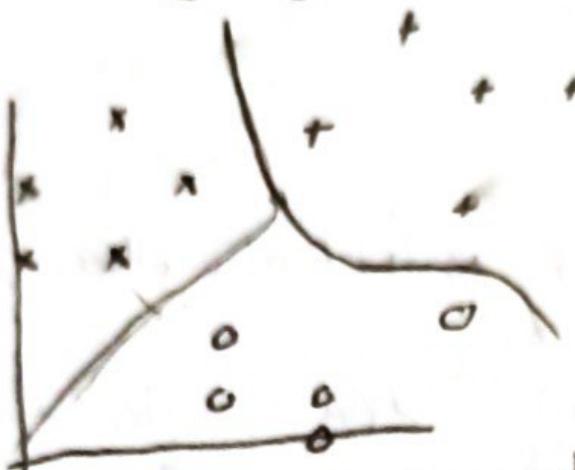
$$b = -0,5 (1,01^2 - 0,29^2) \approx -0,48$$

$$w^T x - b = 0,25x_1 - 0,64x_2 - 0,46x_3 - 0,48$$

③

- ④ It is lazy learner because there is no explicit training phase, we just need to store the labelled data.
- ⑤ the training algorithm does not exist.

⑥



Mais au moins...

Note: in the actual graph from the exam,
(4,4) belongs to w_1 (misclassification)

- ⑦ (According to the actual exam's figure)

$K=1 \quad s_1 \quad s_2 \quad s_3 \quad 66\% \text{ accuracy}$

predict:	3	2	3
	x	✓	✓

$K=3 \quad s_1 \quad s_2 \quad s_3 \quad 66\% \text{ accuracy}$

	1	1	3
✓	x	✓	

(4)

② Sep margin γ is given by $\gamma = \frac{2}{\|w\|}$

$$\gamma = \frac{2}{\|w_3\|} = 2 / \sqrt{1^2 + 1^2} \approx 1.22$$

(b)

One-vs-All:

3 classes
Classifier i is trained w/ i as (+1) and others as (-1)

$$M = \begin{pmatrix} +1 & -1 & -1 \\ -1 & +1 & -1 \\ -1 & -1 & +1 \end{pmatrix}$$

For example, when using SVM 3, w_3 is (+1) while w_1 and w_2 are (-1)

rows \Rightarrow classes
cols \Rightarrow classifiers

One-vs-One:

we need $\frac{k(k-1)}{2} = \frac{3(2)}{2} = 3$ classifiers ✓

Ignored classes are marked as 0

$$M = \begin{pmatrix} +1 & +1 & 0 \\ -1 & 0 & +1 \\ 0 & -1 & -1 \end{pmatrix}$$

⑥ A sample is a support vector if its ^{exactly} on the margin boundary of the canonical hyperplane for at least one classifier (if $|\mathbf{w}^T \mathbf{x} + b| = 1$)

Test SVM1:

$$g_1(\mathbf{x}) = \mathbf{w}_1^T \mathbf{x} + b =$$

$$= \begin{pmatrix} -3 \\ 1 \end{pmatrix} \begin{pmatrix} 3 & 1 \end{pmatrix} + 7 = -9 + 1 + 7 = -1$$

$|-1| = 1 \Rightarrow \mathbf{x}$ is support vector
for SVM1
(on the negative side)

\mathbf{x} is a support vector, now let's check its class

$$g_1(\mathbf{x}) = -1 \rightarrow \text{Sign} = -1$$

$$g_2(\mathbf{x}) = 0,4 \cdot 3 - 1,6 + 1,4 = 1 \rightarrow \text{Sign} = +1$$

$$g_3(\mathbf{x}) = 3 + 1,3 - 6,3 = -4 \rightarrow \text{Sign} = -1$$

so, we have $(-1 \ 1 \ -1)$, which is
the same as row 2 from the OVA matrix
 $\Rightarrow \mathbf{x}$ belongs to ω_2

\downarrow
(Hamming distance = 0)

④

$$g_1(x) = -3 \rightarrow \text{Sign} = -1$$

$$g_2(x) = -0.2 \rightarrow \text{Sign} = -1$$

$$g_3(x) = -1.2 \rightarrow \text{Sign} = -1$$

Yes, this is a special case, the Hamming distance is the same between $(-1 -1 -1)$ and all of the rows in the OVA matrix.

\Rightarrow Ambiguous region

⑤

```
import numpy as np
```

```
def bayes_training(Xtr, Ttr)
```

$X_{-c1} = Xtr[:, Ttr == 1]$ } Get samples from
 $X_{-c2} = Xtr[:, Ttr == 2]$ } each class

$total = Xtr.shape[1]$ } Get priors:

$prior1 = X_{-c1}.shape[1] / total$ } $P(1)$ and $P(1)$

$prior2 = X_{-c2}.shape[1] / total$ } $P(2)$

$mean1 = np.mean(X_{-c1}, axis=1, keepdims=True)$ } Get class

$mean2 = np.mean(X_{-c2}, axis=1, keepdims=True)$ } centroids

$cov1 = np.cov(X_{-c1}, bias=True)$ } Get cov

$cov2 = np.cov(X_{-c2}, bias=True)$ } matrices

(bias=True normalizes by N)

model = }

"mu1": mean1,

"mu2": mean2,

"cov1": cov1,

"cov2": cov2,

"p1": prior1,

"p2": prior2

\searrow return model

```

def gaussian_pdf(X, mu, cov):
    D = X.shape[0]
    X_centered = X - mu
    det_cov = np.linalg.det(cov)
    inv_cov = np.linalg.inv(cov)
    const = 1 / ((2 * np.pi) ** (D / 2) * (det_cov ** 0.5))
    #  $\frac{1}{2\pi^{(D/2)} \sqrt{\det(\text{cov})}}$ 
    exp = -0.5 * np.sum(np.dot(X_centered.T, inv_cov) *
                         X_centered, T, axis=1)
    return const * np.exp(exp)

def bayes_testing(Xte, Tte, model):
    likelihood1 = gaussian_pdf(Xte, model['mu1'], model['cov1'])
    likelihood2 = gaussian_pdf(Xte, model['mu2'], model['cov2'])

    post1 = likelihood1 * model['p1']  $\uparrow P(x|w)$ 
    post2 = likelihood2 * model['p2']  $\left\{ \begin{array}{l} \text{cat posteriors} \\ (\alpha p(x|w_1)p(w_1)) \end{array} \right.$ 
    out = np.where(post1 >= post2, 1, 2)  $\left\{ \begin{array}{l} \text{if } post1 \geq post2 \\ \text{assign 1, else 2} \end{array} \right.$ 
    return out

```

maybe also
 return
 performance
 metrics